

To the Editors

Frontiers in Artificial Intelligence — AI in Finance

Re: Original Research submission — “The AI Incident Disclosure Shadow: An Incident-Level Study Matching AIID-Recorded AI Failures to U.S.-Listed Companies’ Investor Disclosures, 2019–2026.”

Dear Editors,

We submit the above manuscript for consideration as an Original Research article in the AI in Finance specialty section of *Frontiers in Artificial Intelligence*.

The paper studies how realized AI failures at public companies enter — or fail to enter — the securities-disclosure channel, connecting the AI-governance and capital-market-disclosure literatures the journal actively publishes, alongside recent Frontiers in AI work on AI adoption and corporate outcomes. To our knowledge it is the first incident-level match of AI Incident Database (AIID) failures to the SEC filings of the U.S.-listed companies involved. Across 307 matched incidents (21 issuers, 2019–2026), incident-specific disclosure is 1.3% and substantive disclosure 2.3%; the remaining 97.7% form a disclosure “shadow” dominated by substitution of specific events by generic AI risk-factor language, robust to issuer concentration, search window, right-censoring, and an independent full-filing audit.

Every reported statistic is generated by deterministic, seed-fixed scripts, independently re-derived by two verifiers and a manuscript-wide numerical-consistency audit, from hashed pre-registration and SHA-256-pinned inputs. Rare-event inference uses exact Clopper–Pearson intervals and Monte-Carlo permutation tests, and inter-coder reliability is reported.

All data and code are openly available to reviewers now — no account, no request, no waiting — in a published, deterministic Code Ocean capsule (MIT and CC0) at <https://doi.org/10.24433/CO.2340354.v2>. It opens in the browser (no login), showing a file tree on the left with three folders — /data, /code, /results:

- **The dataset is in the /data folder.** The primary file is `data/disclosure_coding.csv` — the 307 matched incidents, one row each, with the four-tier disclosure code and severity. Beside it: `incident_firm_map.csv` (incident-to-issuer matches), `validation_sheet.csv` (human validation), `coding_logs/` (per-incident EDGAR search records, both independent coding passes, and the disagreement/correction logs), and `raw_extracted/` (the AIID incident fields and CSET severity).
- **Pre-computed results are in /results** — every reported number, table, and figure, viewable at a glance without running anything.
- **The code is in /code.** To regenerate every result end to end, click the blue “**Reproducible Run**” button: the pipeline runs offline and deterministically (fixed seed).

The study reports aggregate rates and measures disclosure behavior, not legal compliance; it makes no claim that any named issuer violated a disclosure obligation. It uses only public data (AIID and SEC EDGAR) and involved no human or animal subjects. Use of AI assistance is disclosed in the Methods, the Acknowledgments, and a Generative AI Statement.

The manuscript is original, not under consideration elsewhere, and all authors approve its submission. We have no competing interests to declare.

Sincerely,

Robin Chawla (*corresponding author*)

robin.chawla.cse14@iitbhu.ac.in ORCID 0009-0007-2807-3948

Samir Chincholikar ORCID 0009-0007-2779-3492

on behalf of both authors